deeplearning.ai

Case Studies

Why look at case studies?

Outline

Classic networks:

- LeNet-5
- AlexNet
- VGG

ResNet (152)

Inception

deeplearning.ai

Case Studies

Classic networks

LeNet - 5

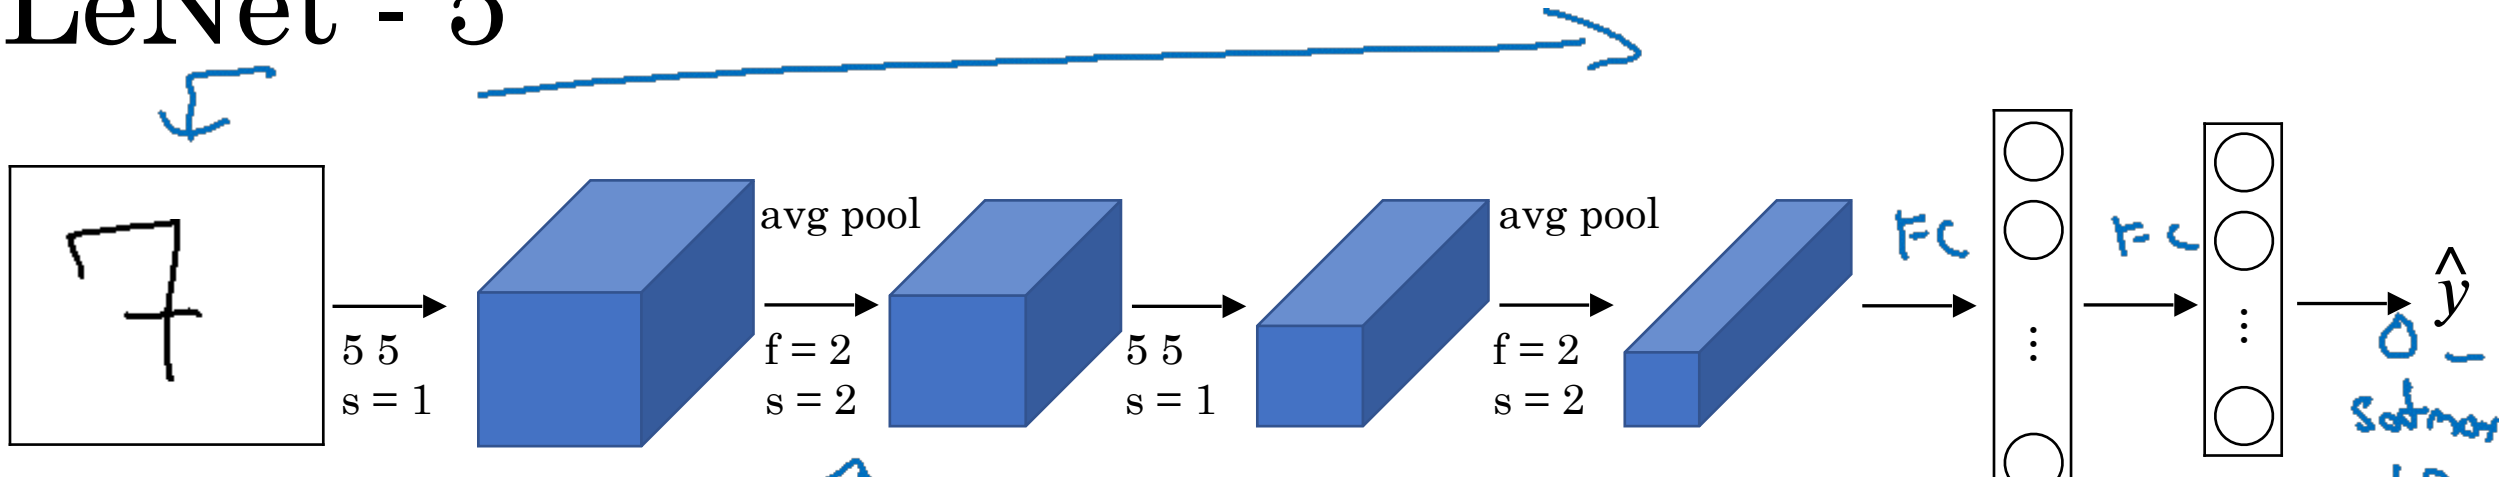

60K parameters.

$n_H, n_w \downarrow$ $n_c \uparrow$

Conv pool Conv pool fc fc output

Activation: Sigmoid/tanh ReLU

[LeCun et al., 1998. Gradient-based learning applied to document recognition]

AlexNet

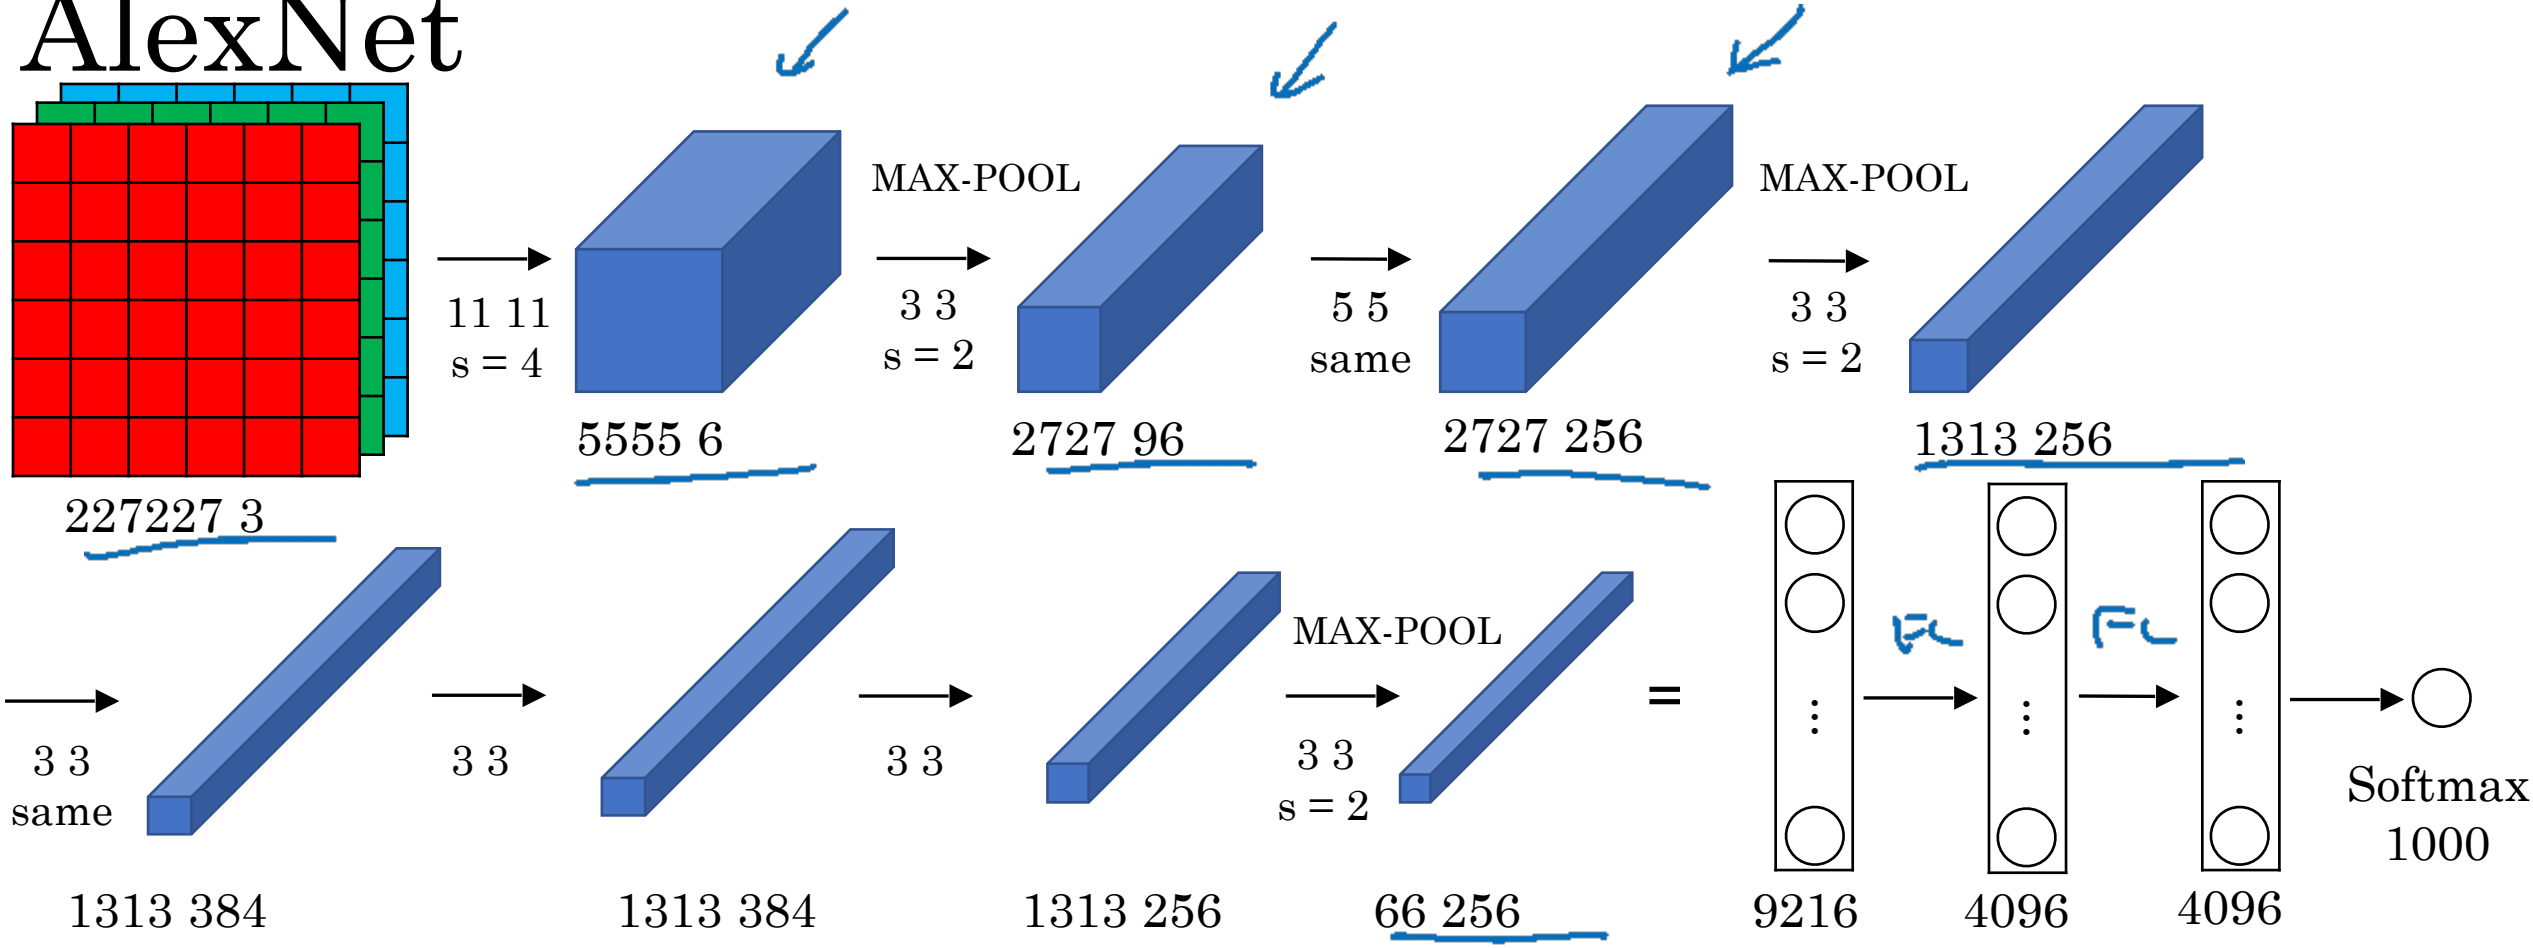

- Similar to LeNet, but much bigger.
- ReLU
- Multiple GPUs.
- Local Response Normalization (LRN)

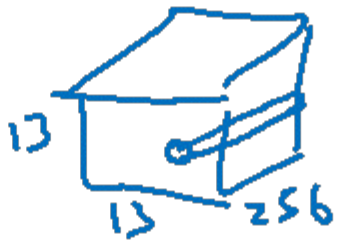

60M parameters

[Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks]

VGG - 16

CONV = 33 filter, s = 1, same MAX-POOL = 22 , s = 2

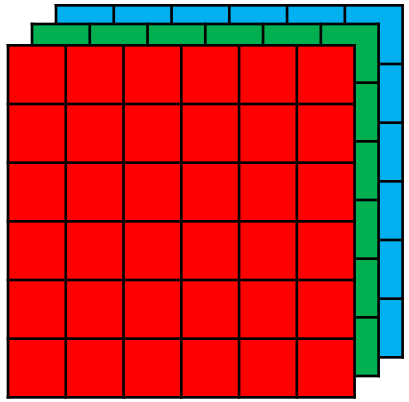

224224 3

→ 22422464 [CONV 64] 2

→ 112112 64 POOL

→ 112112 128 [CONV 128] 2

→ 5656 128 POOL

→ 5656 256 [CONV 256] 3

→ 2828 256 POOL

→ 2828 512 [CONV 512] 3

→ 1414512 POOL

→ 1414 512 [CONV 512] 3

→ 77512 POOL

→ FC 4096

→ FC 4096

→ Softmax 1000

[Simonyan & Zisserman 2015. Very deep convolutional networks for large-scale image recognition]

Andrew Ng

deeplearning.ai

Case Studies

Residual Networks (ResNets)

Residual block

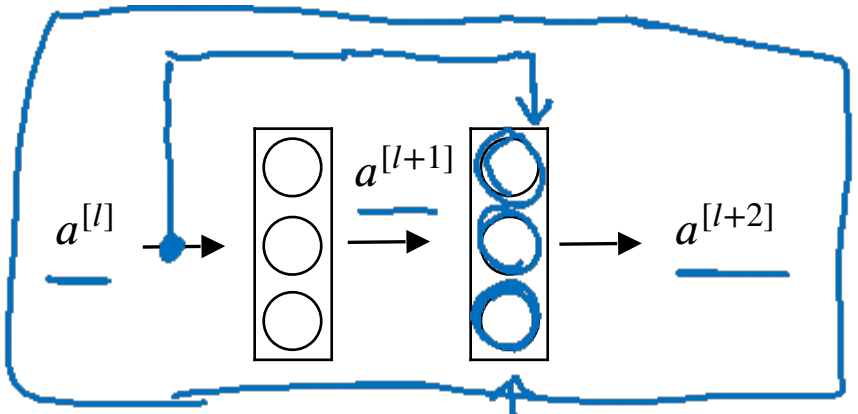

"Short cut" / skip connection

$a^{[l]}$ → Linear → ReLU → $a^{[l+1]}$ → Linear → ReLU → $a^{[l+2]}$

"main path"

$$z^{[l+1]} = W^{[l+1]}a^{[l]} + b^{[l+1]} \qquad a^{[l+1]} = g(z^{[l+1]}) \qquad z^{[l+2]} = W^{[l+2]}a^{[l+1]} + b^{[l+2]} \qquad a^{[l+2]} = g(z^{[l+2]})$$

$$a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$$

[He et al., 2015. Deep residual networks for image recognition]

Residual Network

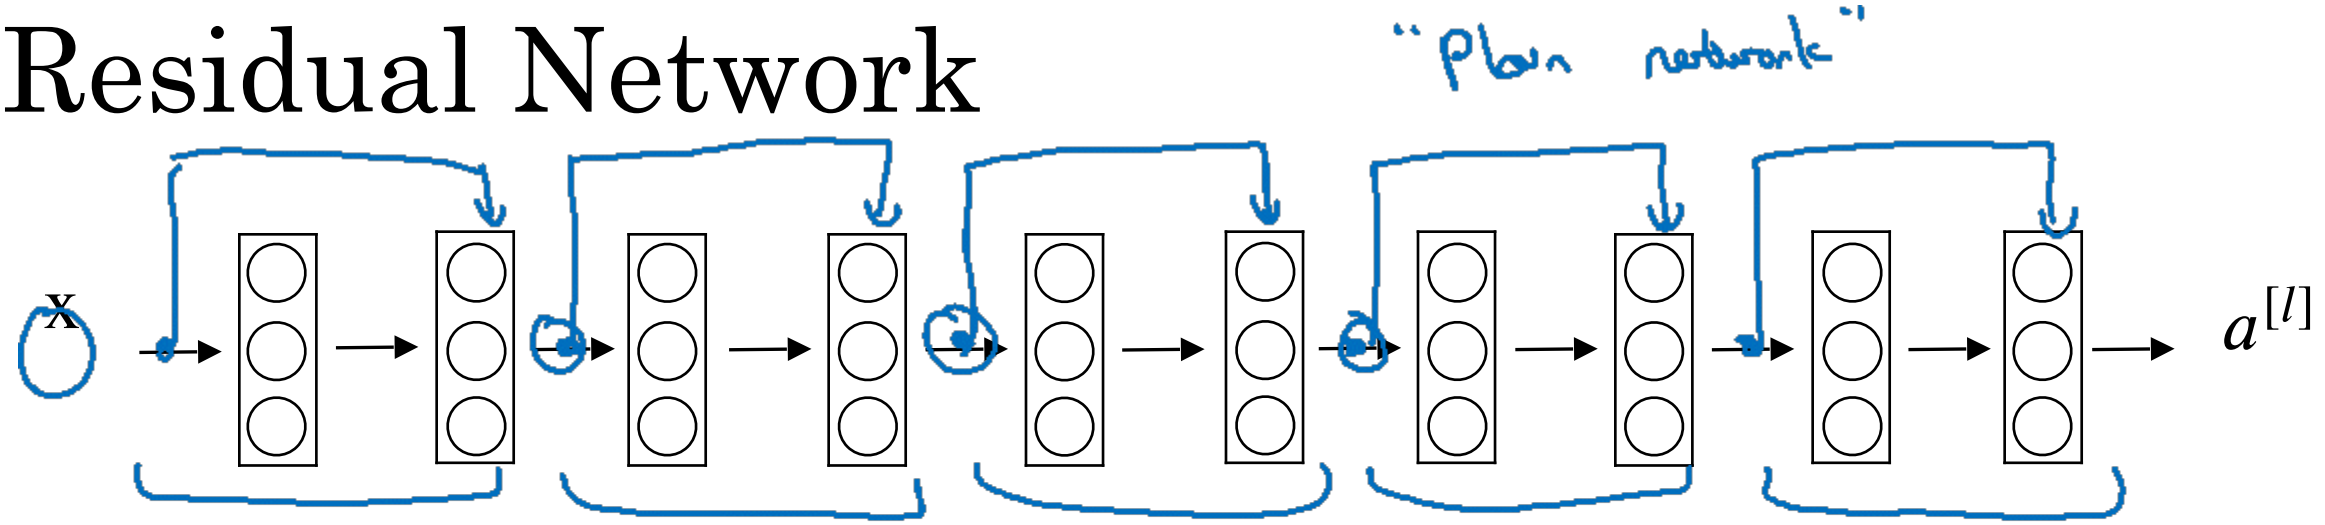

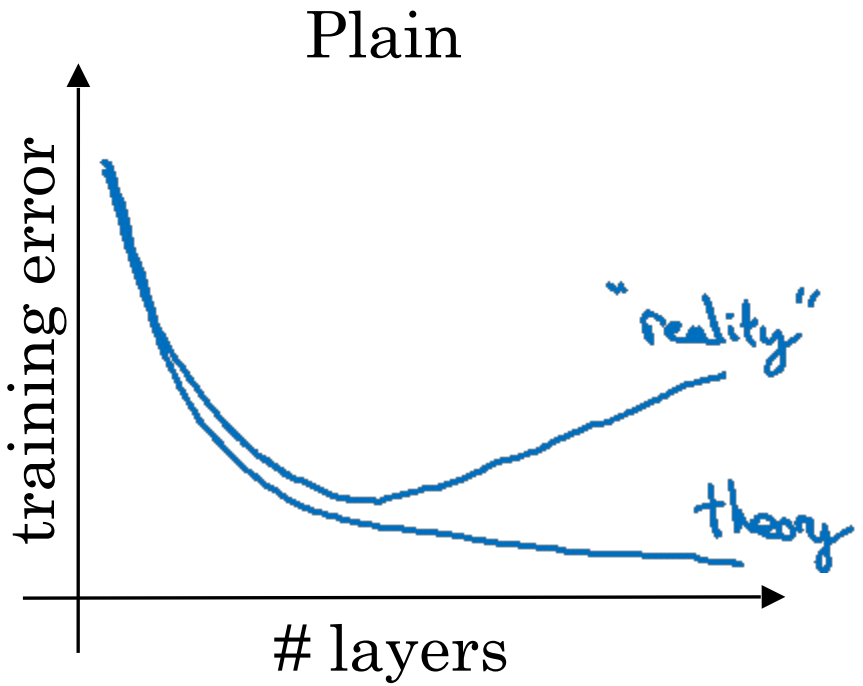

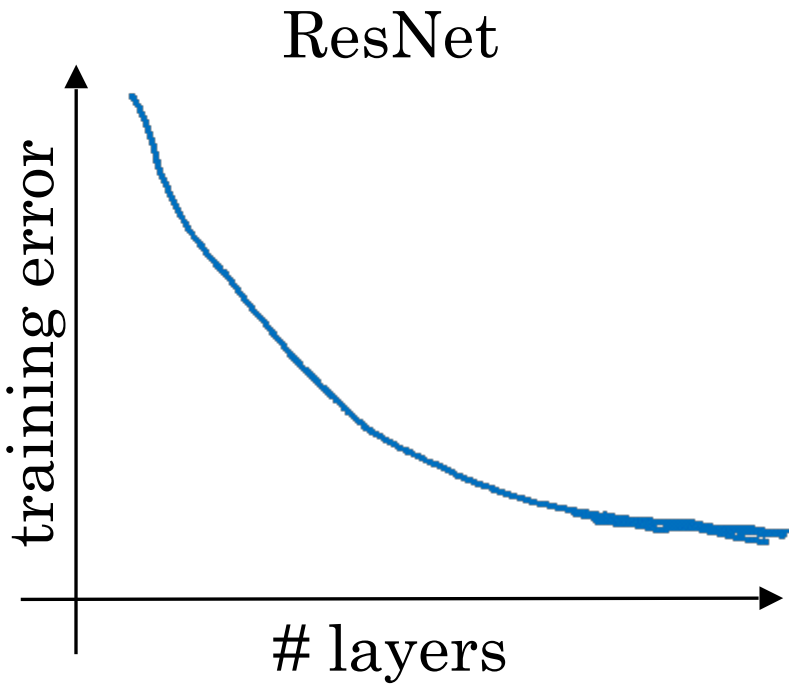

[He et al., 2015. Deep residual networks for image recognition]

deeplearning.ai

Case Studies

Network in Network and 11 convolutions

Why does a 1 × 1 convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

6 6 × 1

*

2

=

2	4	6	...		

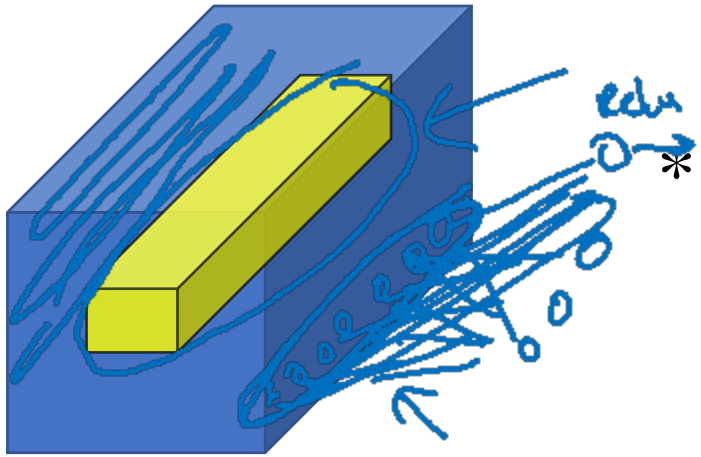

6 6 32

*

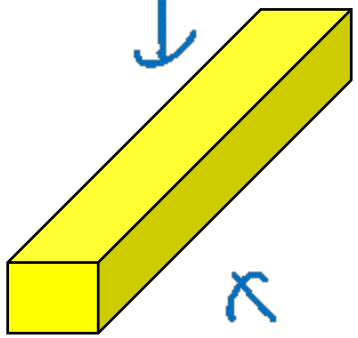

1 1 32

32 → # filters.

$n_c^{[l+1]}$

=

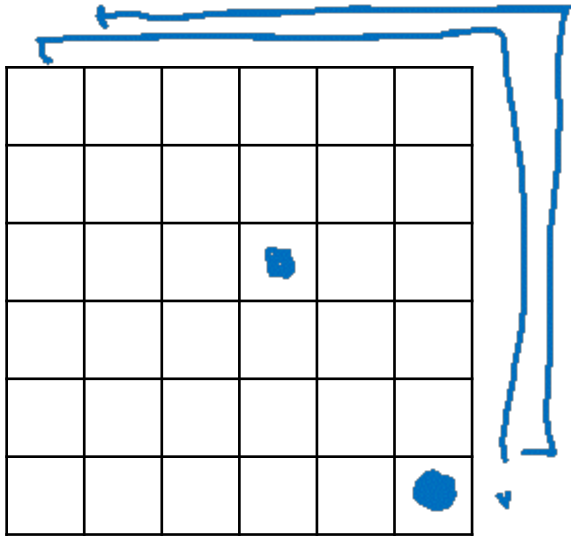

6 6 # filters

ReLU

Network in Network

[Lin et al., 2013. Network in network]

Andrew Ng

Using 11 convolutions

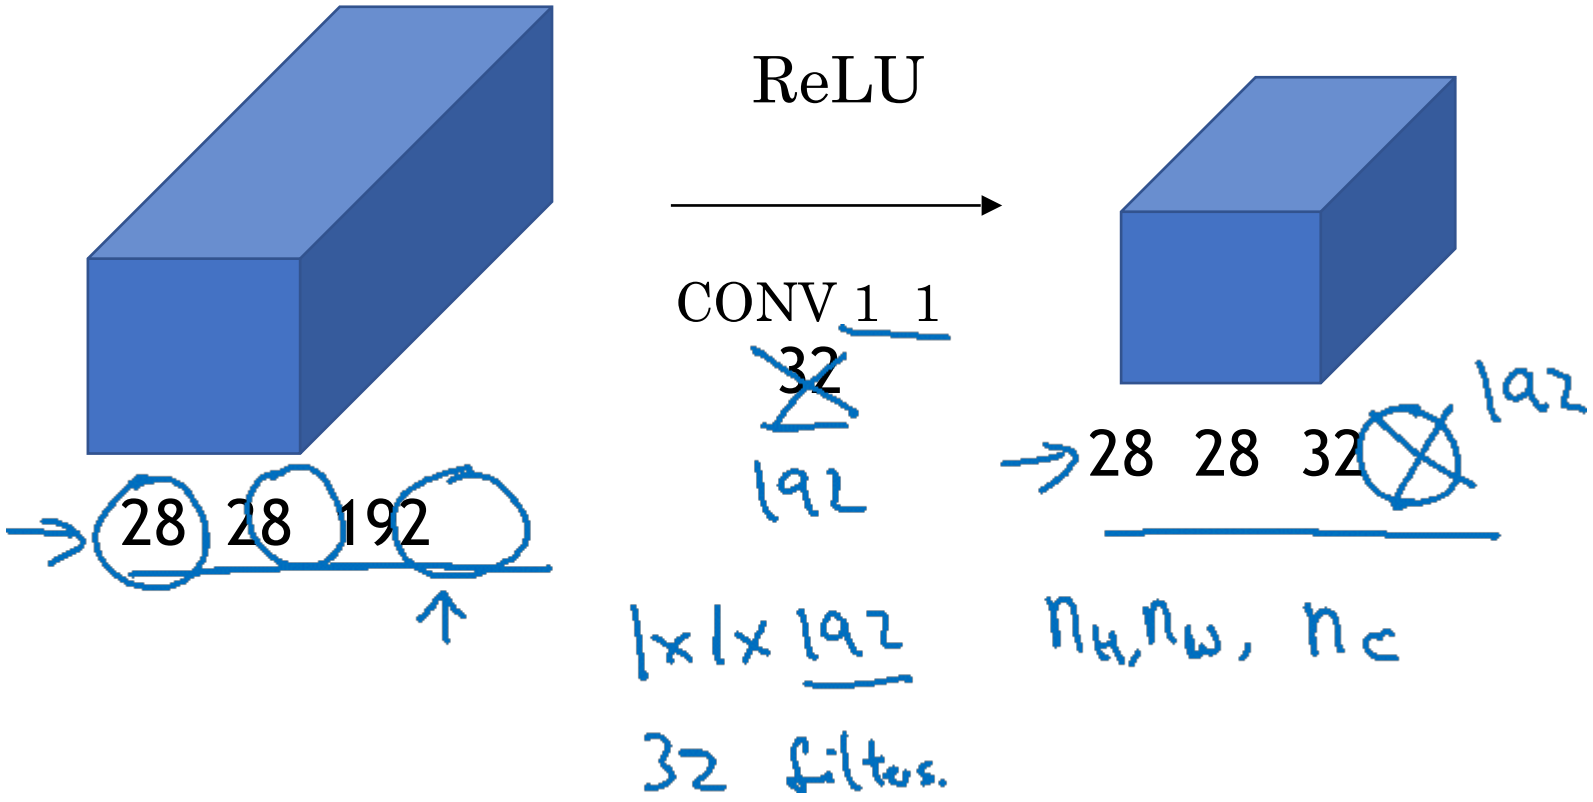

[Lin et al., 2013. Network in network]

deeplearning.ai

Case Studies

Inception network motivation

Motivation for inception network

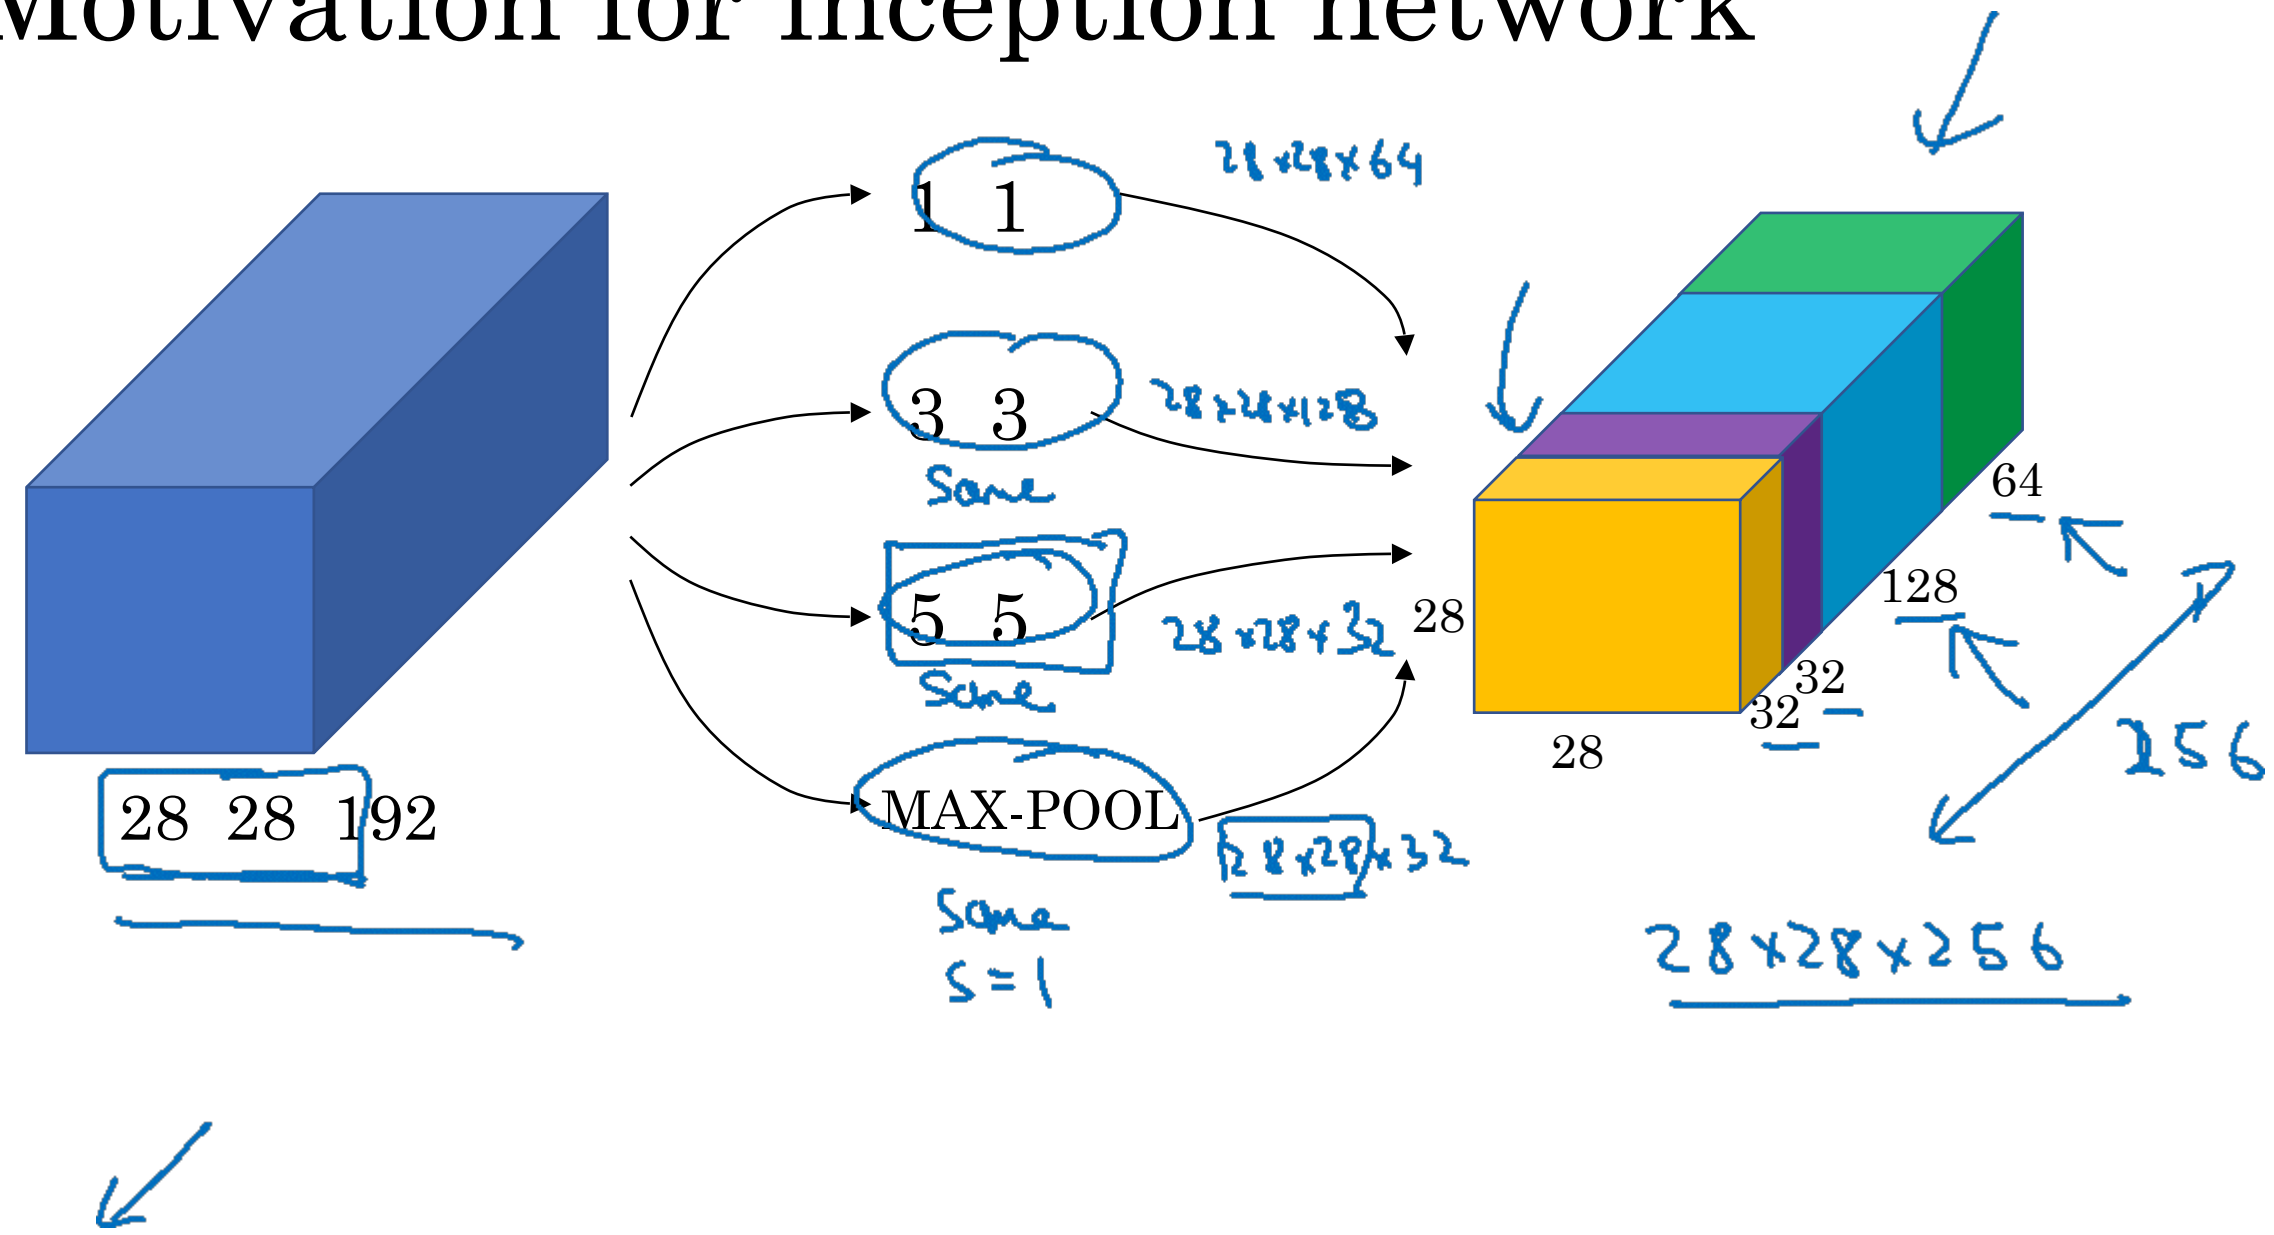

[Szegedy et al. 2014. Going deeper with convolutions]

The problem of computational cost

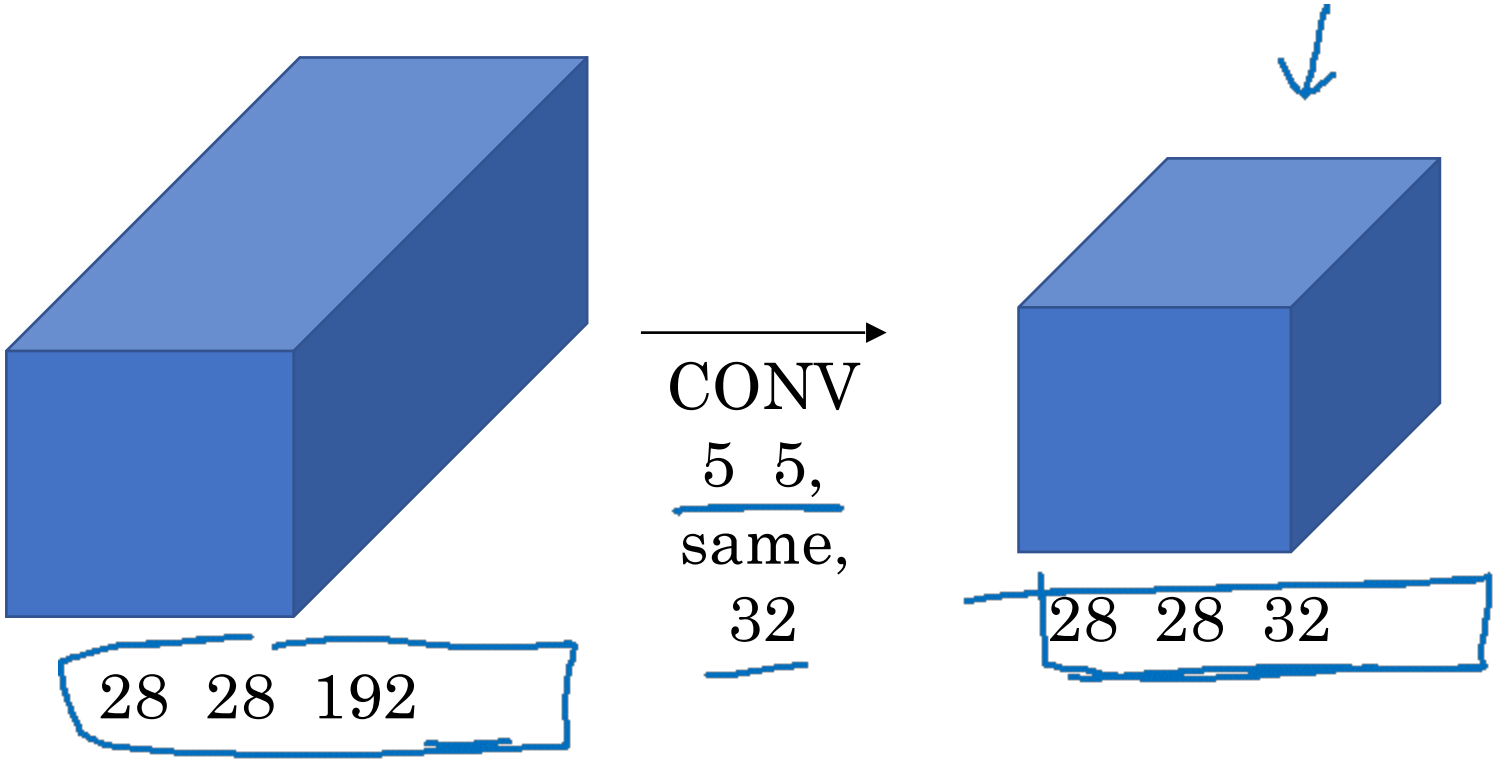

32 filters. filters are 5x5x192.

28x28x32 x 5x5x192 = 120M.

Using 11 convolution

"bottleneck layer"

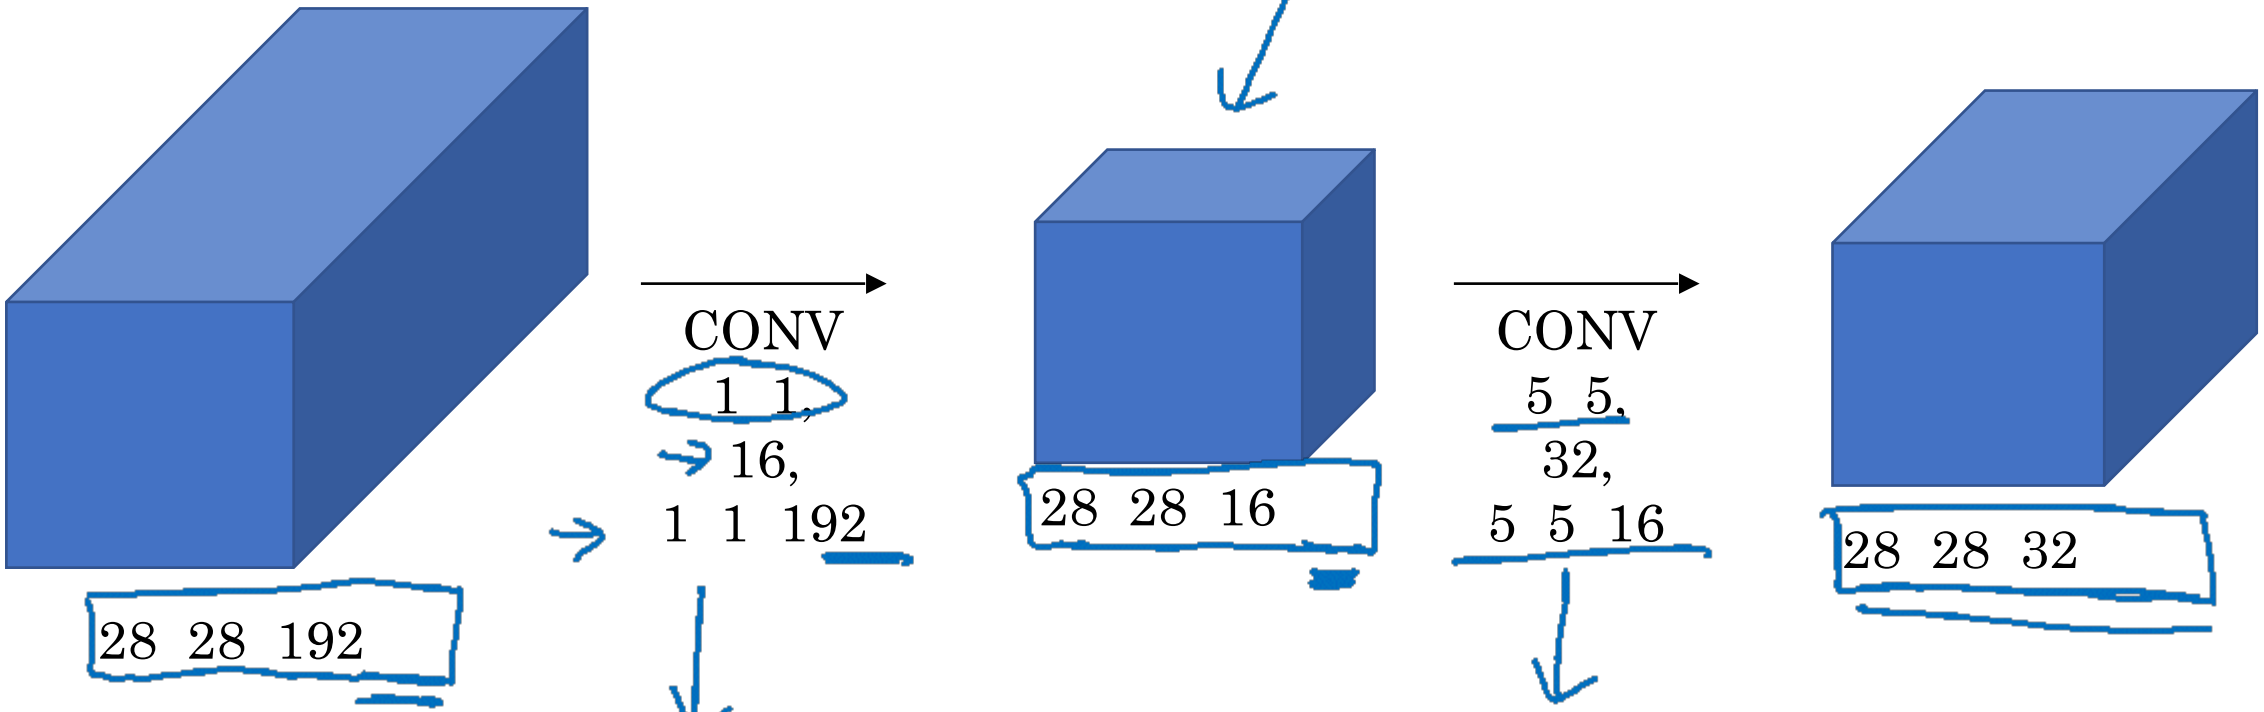

28x28x16 x 192 = 2.4M

28x28x32 x 5x5x16 = 10.0M

12.4M

120M

deeplearning.ai

Case Studies

Inception network

Inception module

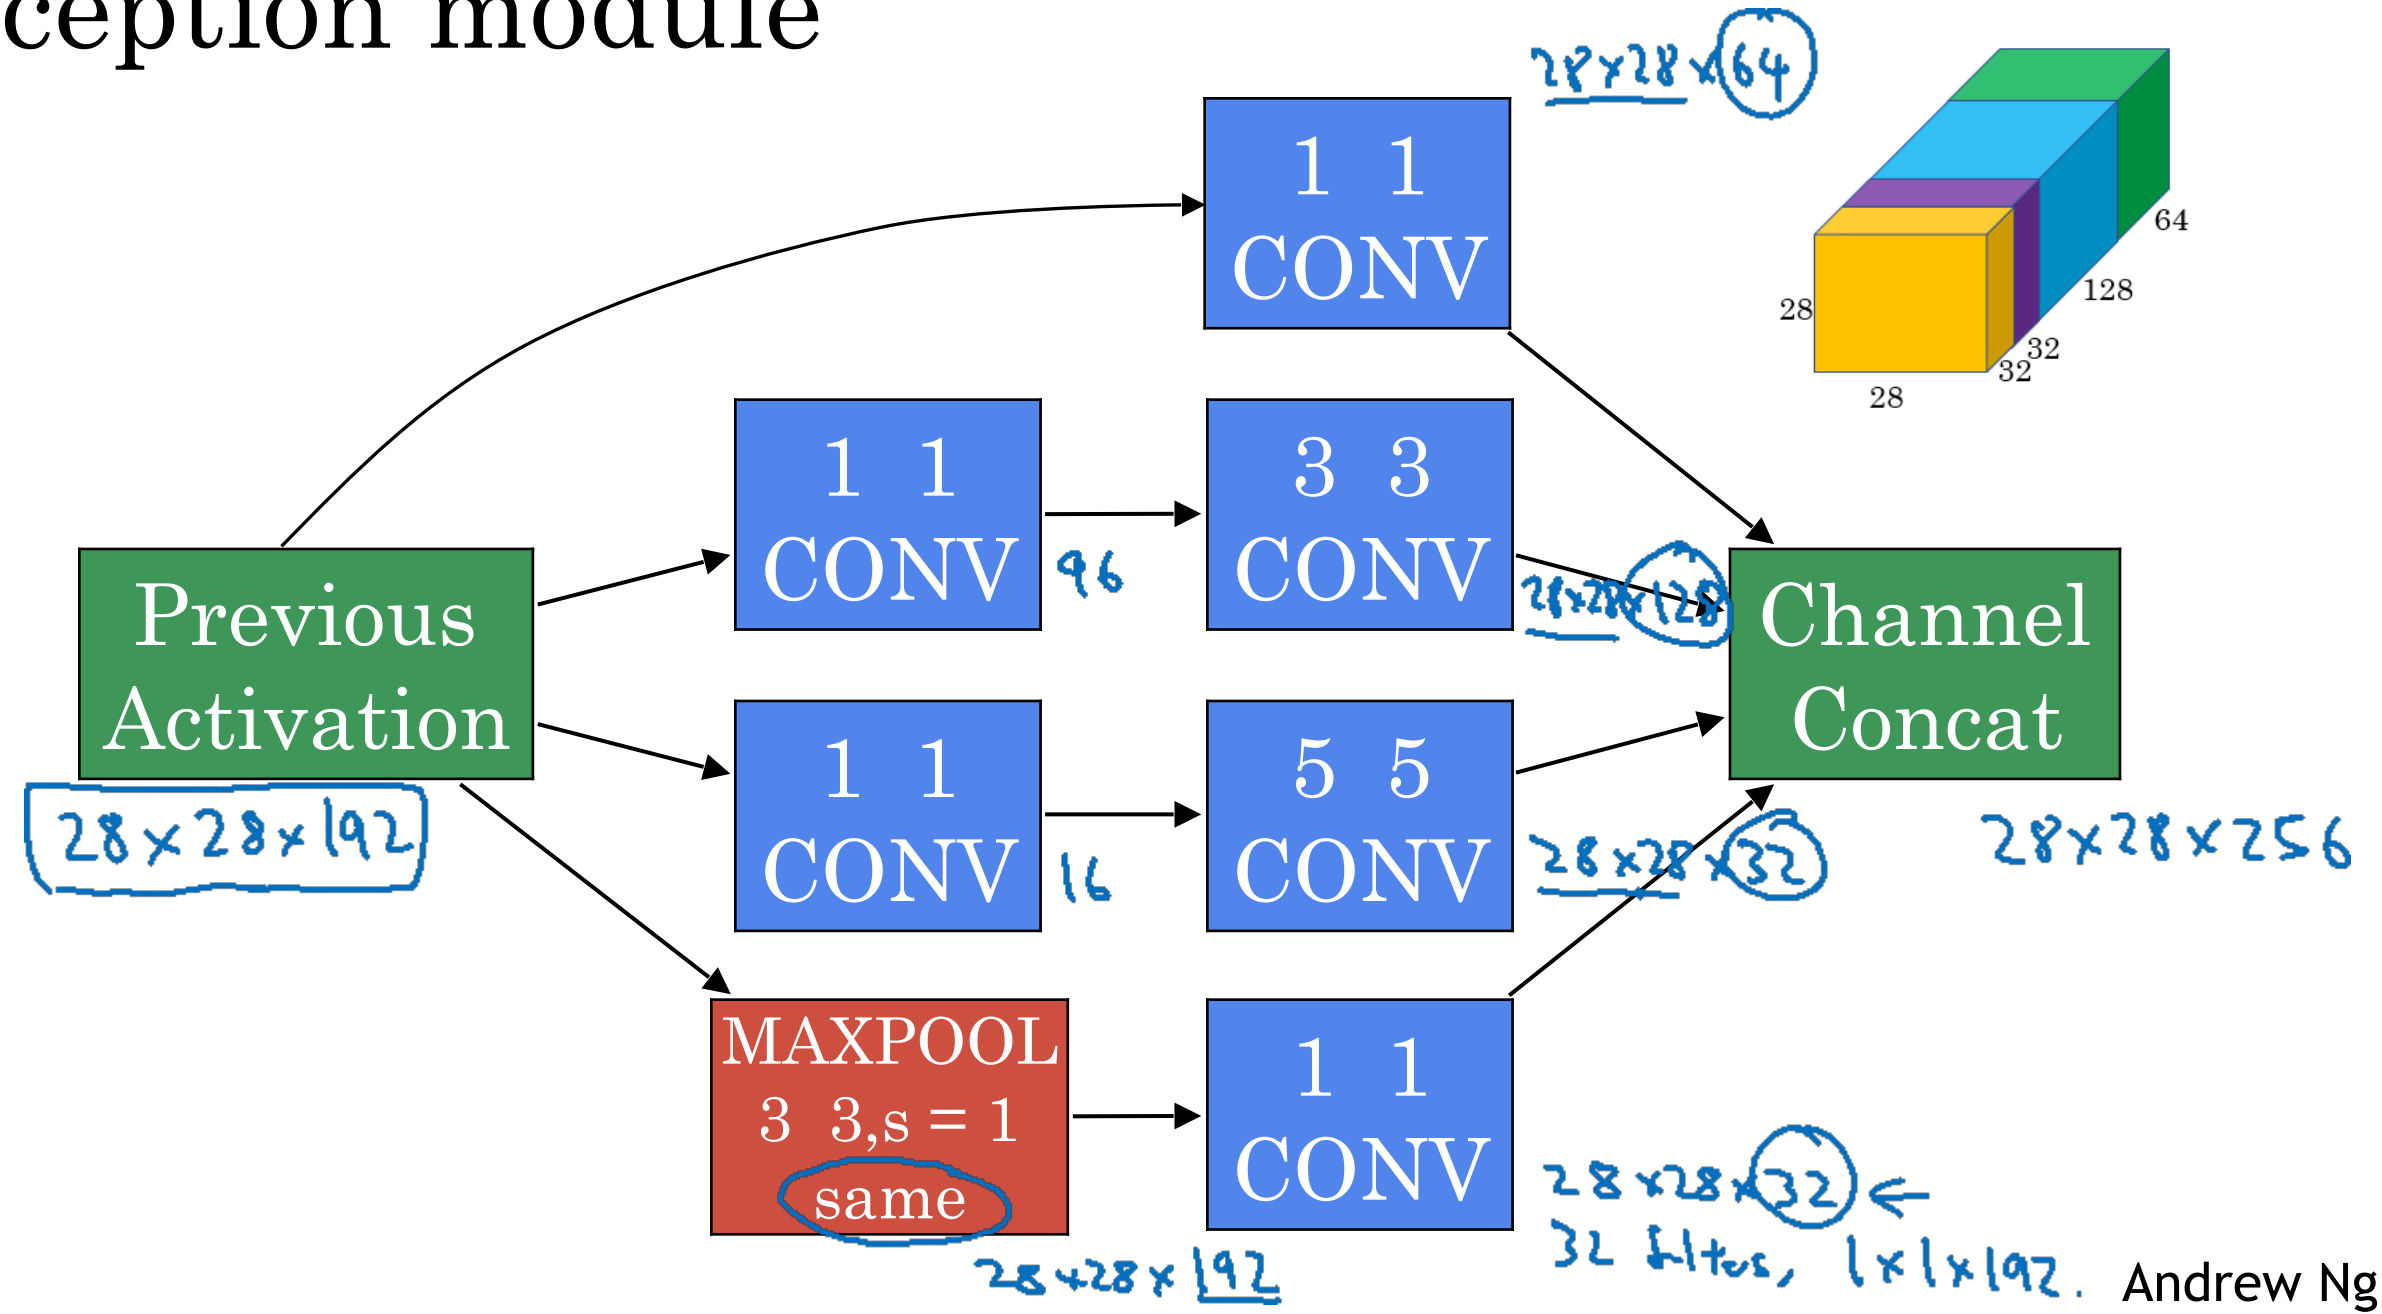

Andrew Ng

Inception network

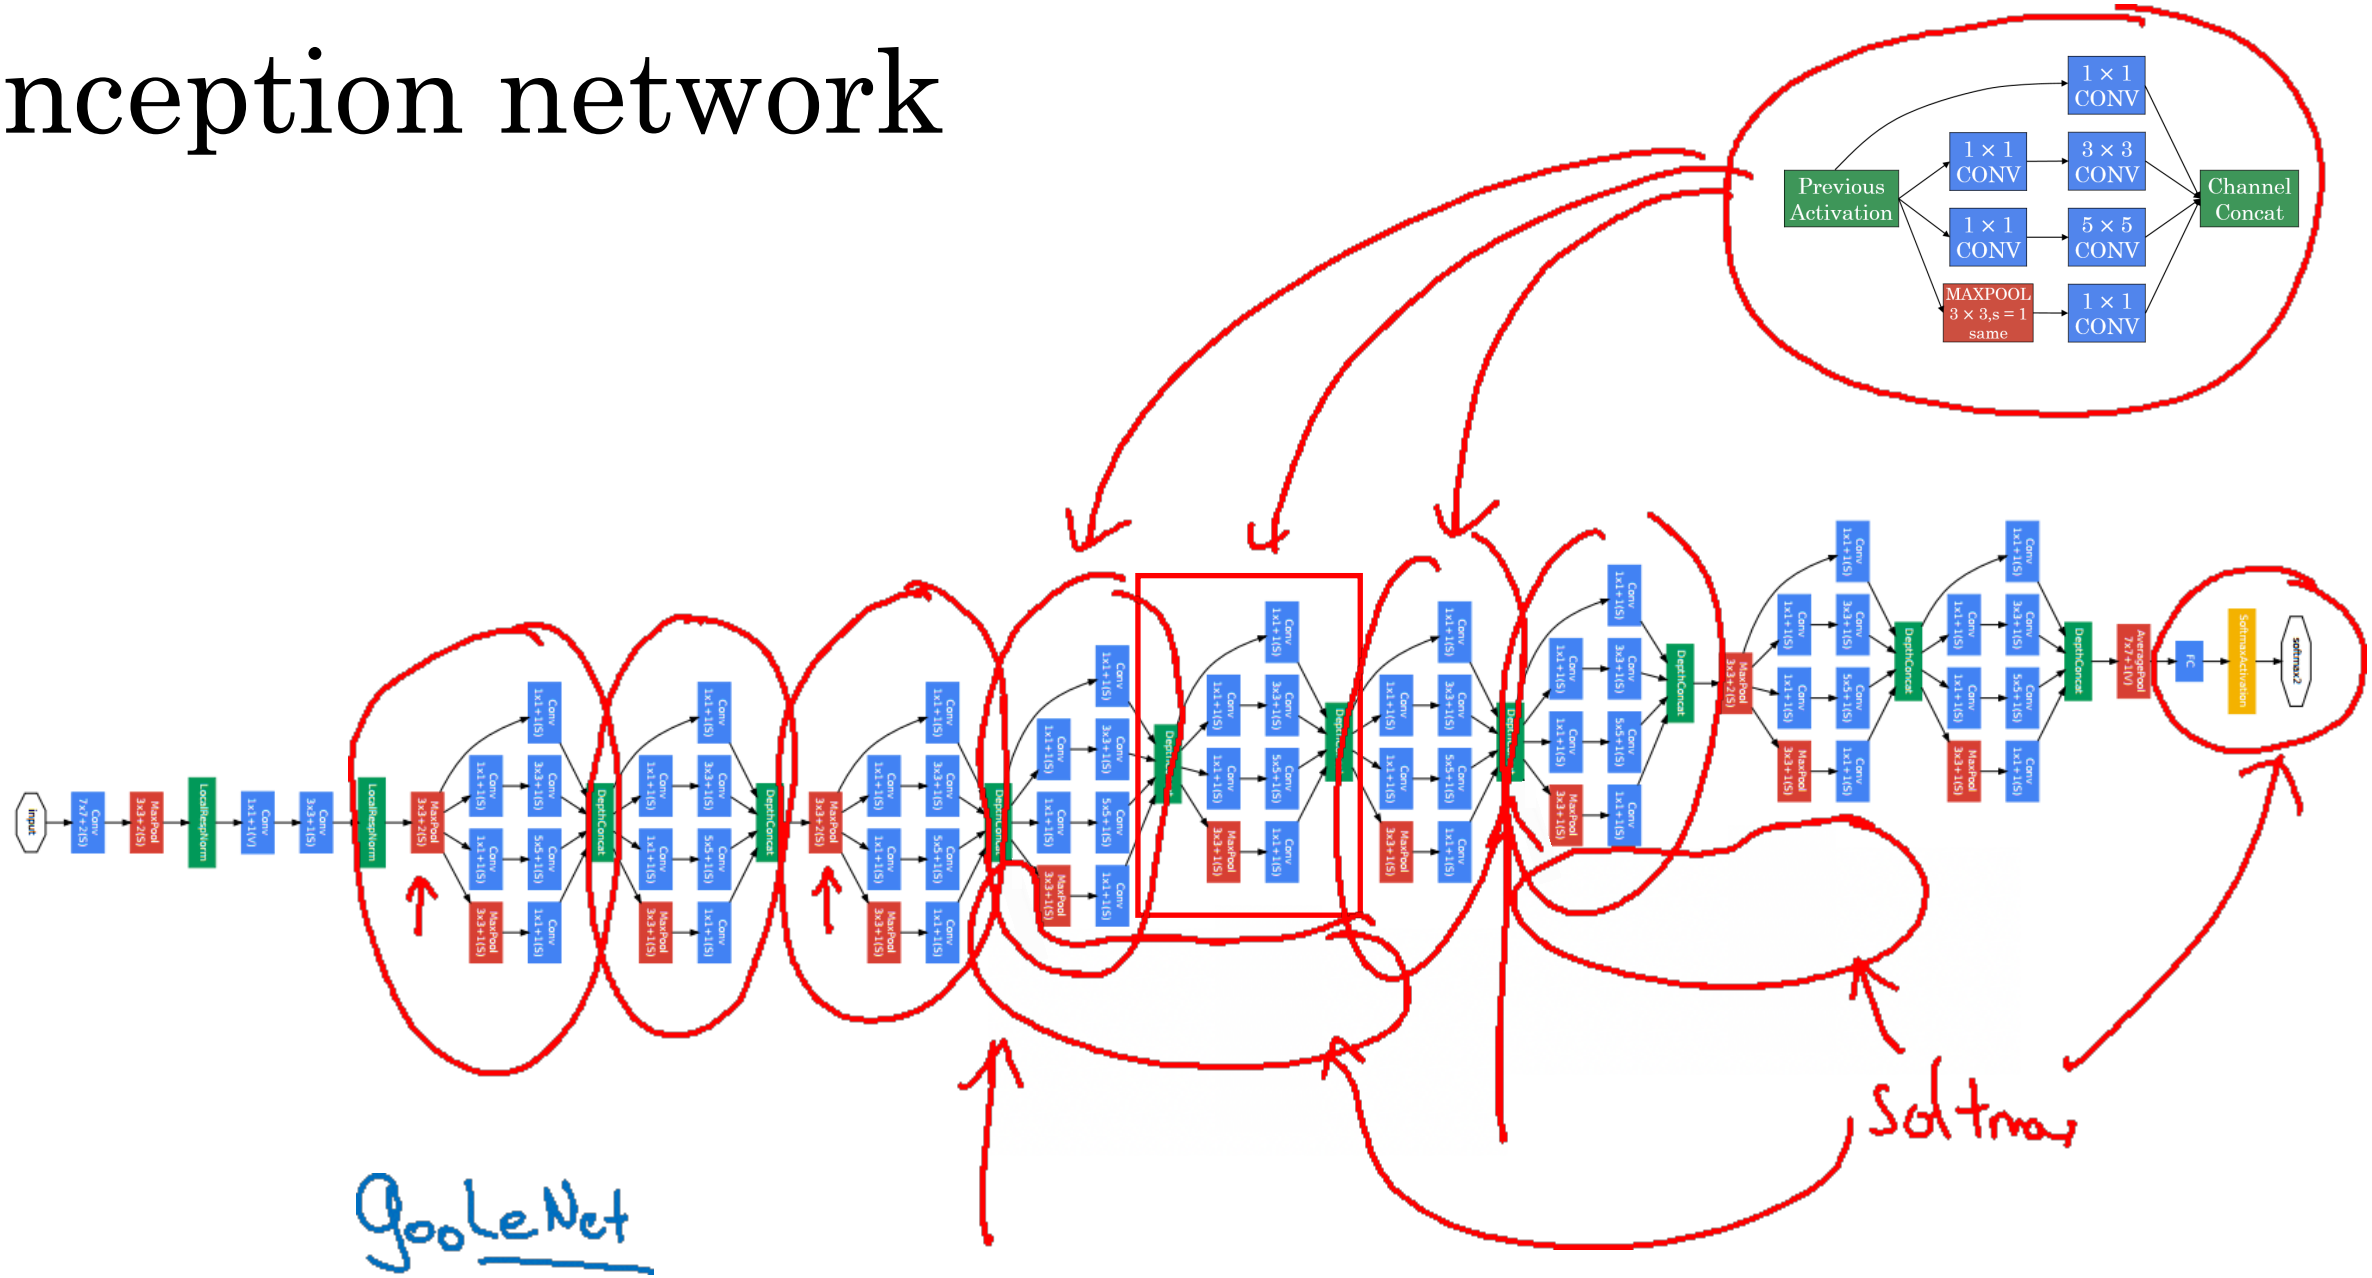

[Szegedy et al., 2014, Going Deeper with Convolutions]

http://knowyourmeme.com/memes/we-need-to-go-deeper